

# Programmazione di Sistemi ~~Embedded e~~ Multicore

Teacher: Daniele De Sensi

Recap

# Recap

- Pi calculation example
- Mutex
- Semaphores
- Barriers and Condition Variables

# Q&A

```
void* fun(void* args){
    int thread_id = *((int*) args);
    printf("I am thread %d\n", thread_id);
    return NULL;
}

int main(){
    ...
    for(int i = 0; i < num_threads; i++){
        pthread_create(&thread_handles[i], NULL, fun, (void*) &i);
    }
    ...
}
```

It does not work, why?

# Q&A

```
void* fun(void* args){  
    int thread_id = *((int*) args);  
    printf("I am thread %d\n", thread_id);  
    return NULL;  
}
```

```
int main(){  
    ...  
    for(int i = 0; i < num_threads; i++){  
        pthread_create(&thread_handles[i], NULL, fun, (void*) &i);  
    }  
    ...  
}
```

It does not work, why?

We have no guarantees that the thread has been created and that it read ((int\*) args) when pthread\_create returns  
&i is a pointer to a variable which might get updated to i+1 before the thread manages to read it

# Q&A

```
void* fun(void* args){
    int thread_id = *((int*) args);
    printf("I am thread %d\n", thread_id);
    return NULL;
}
```

```
int main(){
    ...
    int* ids = (int*) malloc(sizeof(int)*num_thread);
    for(int i = 0; i < num_threads; i++){
        ids[i] = i;
        pthread_create(&thread_handles[i], NULL, fun, (void*) &(ids[i]));
    }
    for(int i = 0; i < num_threads; i++){
        pthread_join(thread_handles[i], NULL);
    }
    free(ids);
    ...
}
```

Correct way of doing it

Questions?


# Read-Write Locks



# Controlling access to a large, shared data structure

- Let's look at an example.
- Suppose the shared data structure is a **sorted** linked list of ints, and the operations of interest are Member, Insert, and Delete.

# Linked Lists

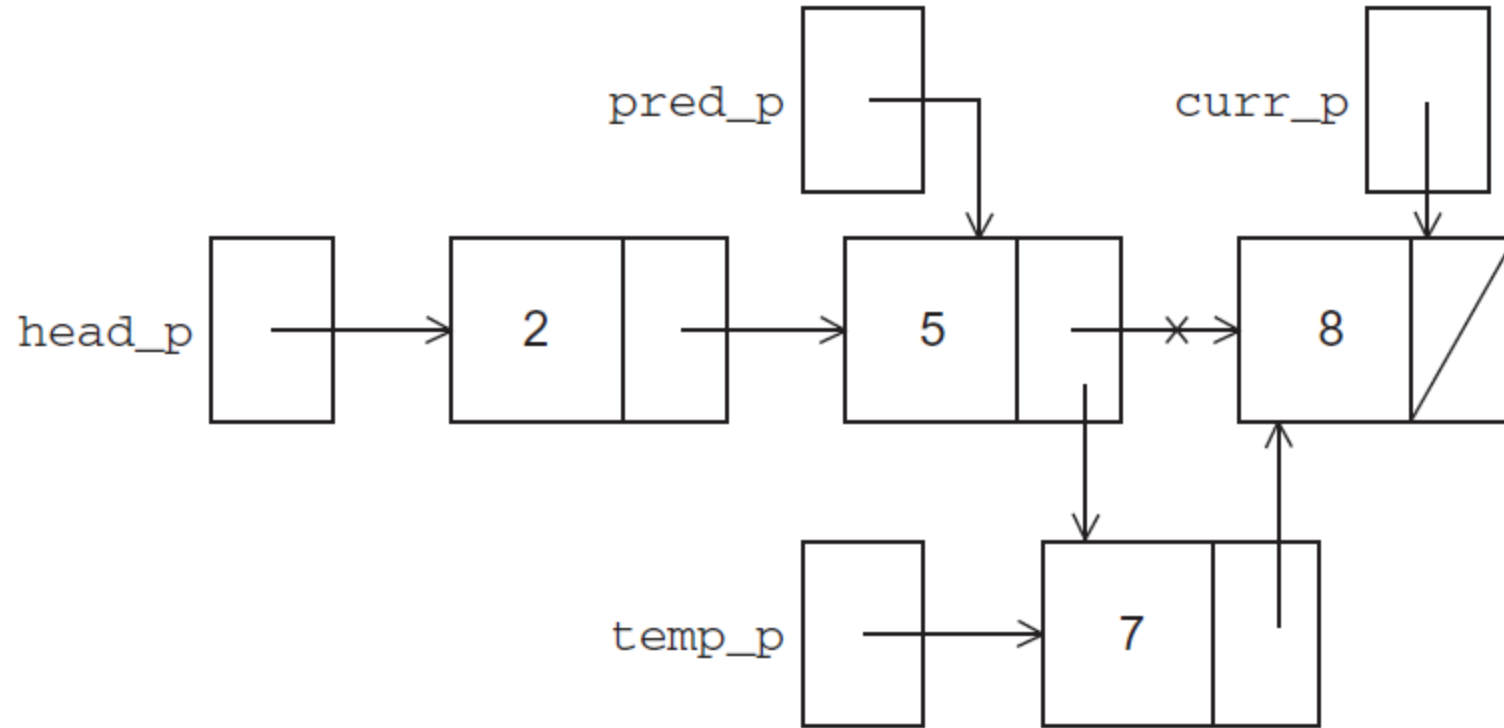


```
struct list_node_s {  
    int data;  
    struct list_node_s* next;  
}
```

# Linked List Membership

```
int Member(int value, struct list_node_s* head_p) {  
    struct list_node_s* curr_p = head_p;  
  
    while (curr_p != NULL && curr_p->data < value)  
        curr_p = curr_p->next;  
  
    if (curr_p == NULL || curr_p->data > value) {  
        return 0;  
    } else {  
        return 1;  
    }  
} /* Member */
```

# Inserting a new node into a list



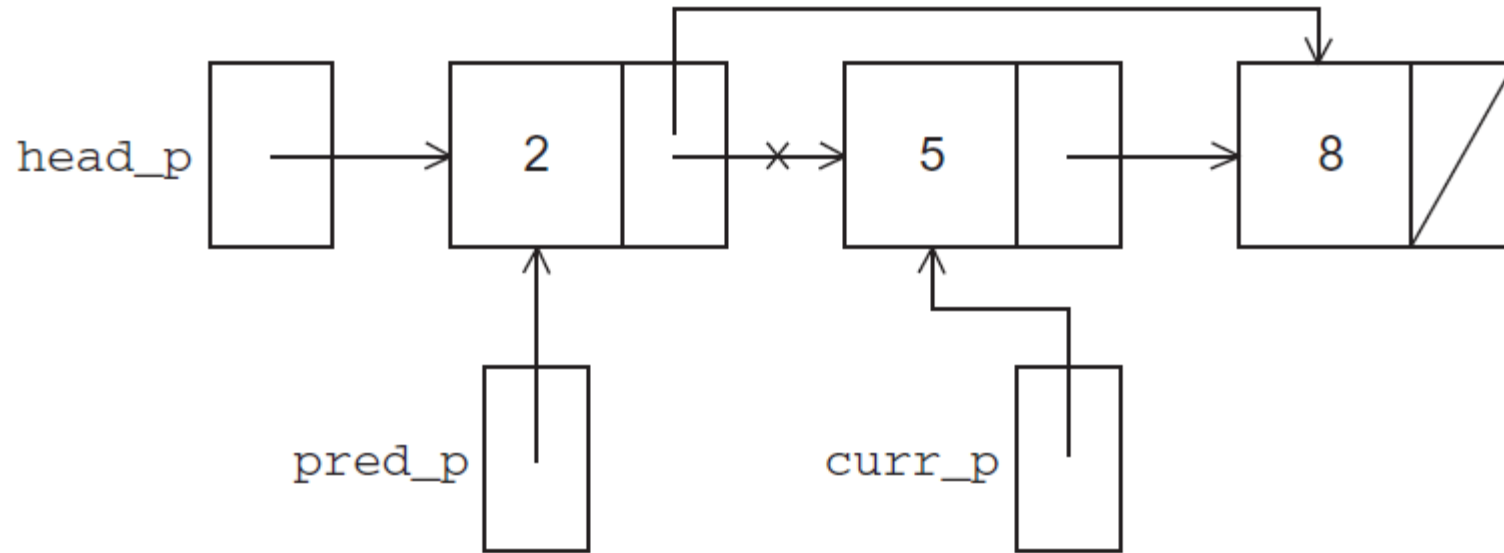
# Inserting a new node into a list

```
int Insert(int value, struct list_node_s** head_pp) {
    struct list_node_s* curr_p = *head_pp;
    struct list_node_s* pred_p = NULL;
    struct list_node_s* temp_p;

    while (curr_p != NULL && curr_p->data < value) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }

    if (curr_p == NULL || curr_p->data > value) {
        temp_p = malloc(sizeof(struct list_node_s));
        temp_p->data = value;
        temp_p->next = curr_p;
        if (pred_p == NULL) /* New first node */
            *head_pp = temp_p;
        else
            pred_p->next = temp_p;
        return 1;
    } else { /* Value already in list */
        return 0;
    }
} /* Insert */
```

# Deleting a node from a linked list



# Deleting a node from a linked list

```
int Delete(int value, struct list_node_s** head_pp) {
    struct list_node_s* curr_p = *head_pp;
    struct list_node_s* pred_p = NULL;

    while (curr_p != NULL && curr_p->data < value) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }

    if (curr_p != NULL && curr_p->data == value) {
        if (pred_p == NULL) { /* Deleting first node in list */
            *head_pp = curr_p->next;
            free(curr_p);
        } else {
            pred_p->next = curr_p->next;
            free(curr_p);
        }
        return 1;
    } else { /* Value isn't in list */
        return 0;
    }
} /* Delete */
```

Questions?

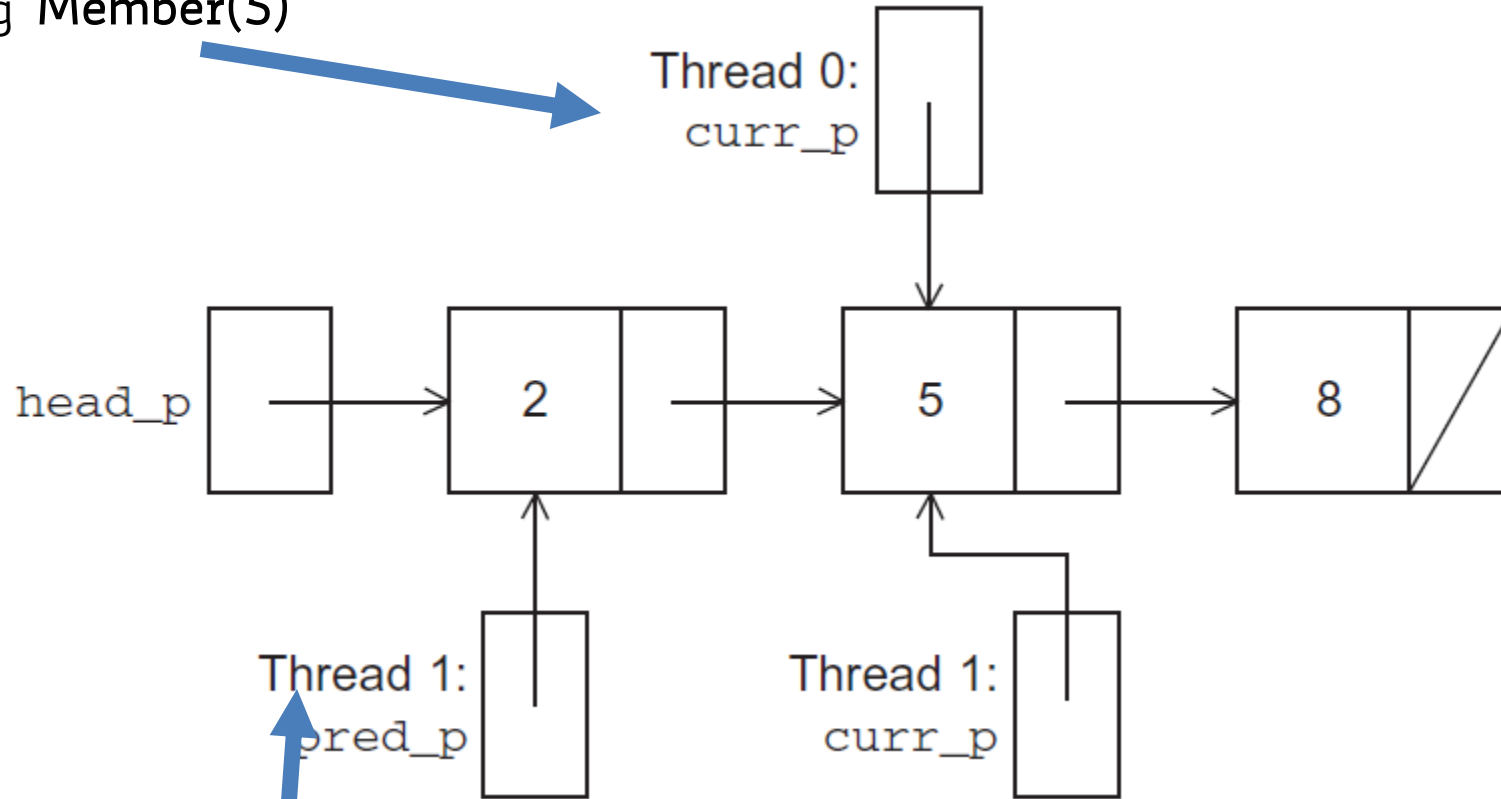


# A Multi-Threaded Linked List

- Let's try to use these functions in a Pthreads program.
- In order to share access to the list, we can define `head_p` to be a global variable.
- This will simplify the function headers for `Member`, `Insert`, and `Delete`, since we won't need to pass in either `head_p` or a pointer to `head_p`: we'll only need to pass in the value of interest.
- If multiple threads call `Member` at the same time, we are fine
- What if one thread calls `Member` while another thread is deleting an element?

# Simultaneous access by two threads

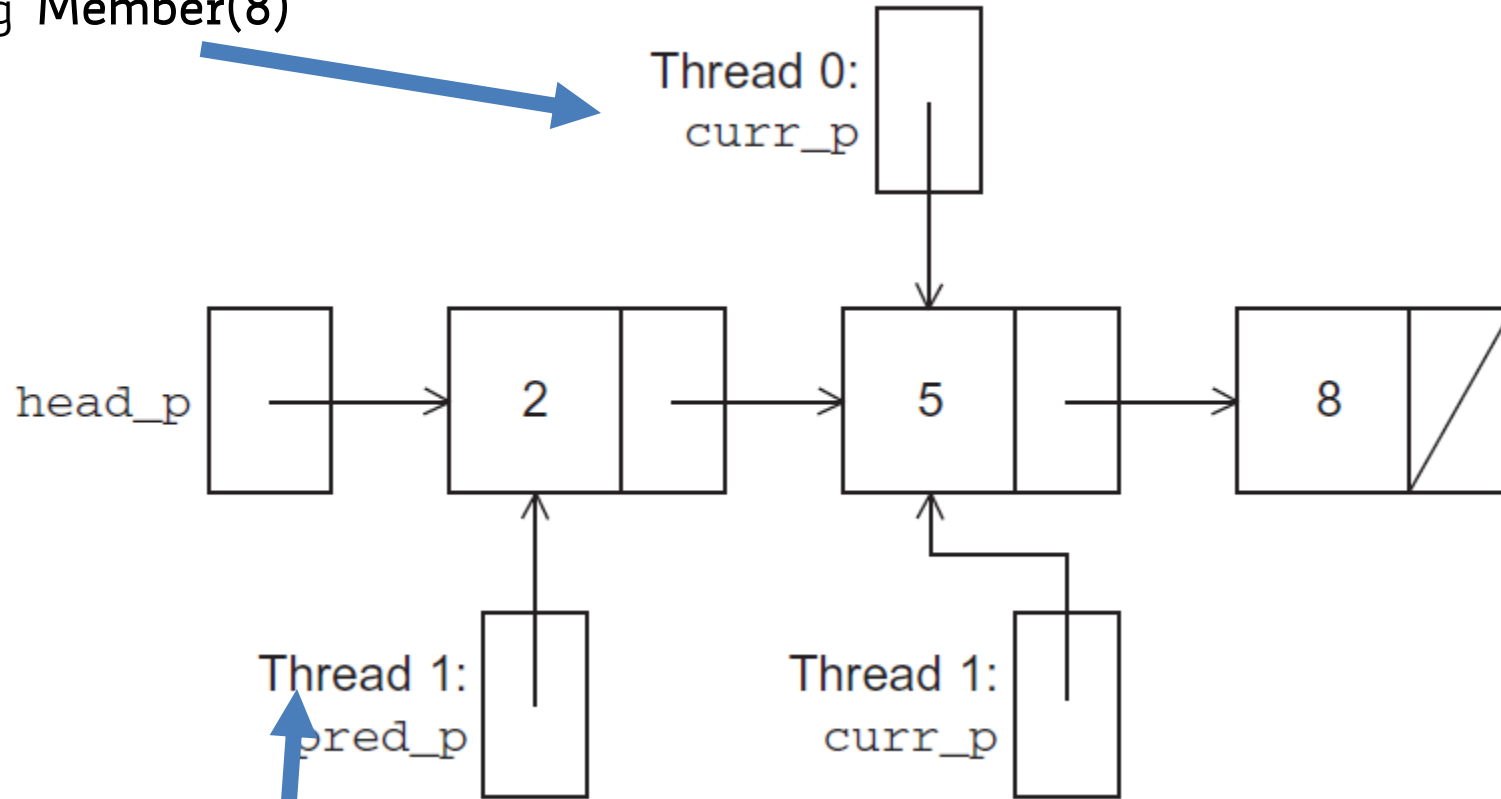
Executing **Member(5)**



Executing **Delete(5)**

# Simultaneous access by two threads

Executing **Member(8)**



Executing **Delete(8)**

# Solution #1

- An obvious solution is to simply lock the list any time that a thread attempts to access it.
- A call to each of the three functions can be protected by a mutex.

```
Pthread_mutex_lock(&list_mutex);  
Member(value);  
Pthread_mutex_unlock(&list_mutex);
```

# Issues

- We're serializing access to the list.
- If the vast majority of our operations are calls to **Member**, we'll fail to exploit this opportunity for parallelism.
- On the other hand, if most of our operations are calls to **Insert** and **Delete**, then this may be the best solution since we'll need to serialize access to the list for most of the operations, and this solution will certainly be easy to implement.

# Solution #2

- Instead of locking the entire list, we could try to lock individual nodes.
- A “finer-grained” approach.

```
struct list_node_s {  
    int data;  
    struct list_node_s* next;  
    pthread_mutex_t mutex;  
}
```

# Issues

- This is much more complex than the original `Member` function.
- It is also much slower, since, in general, each time a node is accessed, a mutex must be locked and unlocked.
- The addition of a mutex field to each node will substantially increase the amount of storage needed for the list.

Questions?



# Pthreads Read-Write Locks

- Neither of our multi-threaded linked lists exploits the potential for simultaneous access to any node by threads that are executing Member.
- The first solution only allows one thread to access the entire list at any instant.
- The second only allows one thread to access any given node at any instant.

# Pthreads Read-Write Locks

- A read-write lock is somewhat like a mutex except that it provides two lock functions.
- The first lock function locks the read-write lock for reading, while the second locks it for writing.

# Pthreads Read-Write Locks

- So multiple threads can simultaneously obtain the lock by calling the read-lock function, while only one thread can obtain the lock by calling the write-lock function.
- Thus, if any threads own the lock for reading, any threads that want to obtain the lock for writing will block in the call to the write-lock function.
- If any thread owns the lock for writing, any threads that want to obtain the lock for reading or writing will block in their respective locking functions.

# pthread\_rwlock functions

pthread\_rwlock\_init initializes the rwlock

```
int pthread_rwlock_init(pthread_rwlock_t* rwlock,  
                        pthread_rwlockattr_t* attr);
```

pthread\_rwlock\_destroy frees the rwlock

```
int pthread_rwlock_destroy(pthread_rwlock_t* rwlock);
```

We can set it to NULL.

Used to specify if readers must have priority over writers, or viceversa.

Check pthread\_rwlockattr\_setkind\_np

# Protecting our linked list functions

```
pthread_rwlock_rdlock(&rwlock);  
Member(value);  
pthread_rwlock_unlock(&rwlock);  
.  
.  
.  
pthread_rwlock_wrlock(&rwlock);  
Insert(value);  
pthread_rwlock_unlock(&rwlock);  
.  
.  
.  
pthread_rwlock_wrlock(&rwlock);  
Delete(value);  
pthread_rwlock_unlock(&rwlock);
```

# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	0.213	0.123	0.098	0.115
One Mutex for Entire List	0.211	0.450	0.385	0.457
One Mutex per Node	1.680	5.700	3.450	2.700

100,000 ops/thread

99.9% Member

0.05% Insert

0.05% Delete

8 (physical) cores

# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	0.213	0.123	0.098	0.115
One Mutex for Entire List	0.211	0.450	0.385	0.457
One Mutex per Node	1.680	5.700	3.450	2.700

100,000 ops/thread

99.9% Member

0.05% Insert

0.05% Delete

8 (physical) cores

# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	0.213	0.123	0.098	0.115
One Mutex for Entire List	0.211	0.450	0.385	0.457
One Mutex per Node	1.680	5.700	3.450	2.700

100,000 ops/thread

99.9% Member

0.05% Insert

0.05% Delete

8 (physical) cores



# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	2.48	4.97	4.69	4.71
One Mutex for Entire List	2.50	5.13	5.04	5.11
One Mutex per Node	12.00	29.60	17.00	12.00

100,000 ops/thread

80% Member

10% Insert

10% Delete

8 (physical) cores

# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	2.48	4.97	4.69	4.71
One Mutex for Entire List	2.50	5.13	5.04	5.11
One Mutex per Node	12.00	29.60	17.00	12.00

100,000 ops/thread

80% Member

10% Insert

10% Delete

8 (physical) cores

# Take-home message

Read-write locks give performance gains as long as the number of insertion/deletion is small compared to the number of member operations

Questions?

# Thread-Safety

# Example

- A block of code is **thread-safe** if it can be simultaneously executed by multiple threads without causing problems.
- Suppose we want to use multiple threads to “tokenize” a file that consists of ordinary English text.
- The tokens are just contiguous sequences of characters separated from the rest of the text by white-space — a space, a tab, or a newline.

# Simple approach

- Divide the input file into lines of text and assign the lines to the threads in a round-robin fashion.
- The first line goes to thread 0, the second goes to thread 1, . . . , the  $t$ th goes to thread  $t$ , the  $t + 1$ st goes to thread 0, etc.

# Simple approach

- We can serialize access to the lines of input using semaphores.
  - Why semaphores and not mutex?
- After a thread has read a single line of input, it can tokenize the line using the `strtok` function.



# The strtok function

- The first time it's called the string argument should be the text to be tokenized.
  - Our line of input.
- For subsequent calls, the first argument should be NULL.

```
char* strtok(  
    char*      string      /* in/out */,  
    const char* separators /* in    */ );
```

# The strtok function

- The idea is that in the first call, `strtok` caches a pointer to string, and for subsequent calls it returns successive tokens taken from the cached pointer.

# Multi-threaded tokenizer (1)

sems[0] initialized to 1

sems[i] (with  $i > 0$ ) initialized to 0

```
void *Tokenize(void* rank) {  
    long my_rank = (long) rank;  
    int count;  
    int next = (my_rank + 1) % thread_count;  
    char *fg_rv;  
    char my_line[MAX];  
    char *my_string;  
  
    sem_wait(&sems[my_rank]);  
    fg_rv = fgets(my_line, MAX, stdin);  
    sem_post(&sems[next]);  
    while (fg_rv != NULL) {  
        printf("Thread %ld > my line = %s", my_rank, my_line);
```

# Multi-threaded tokenizer (2)

```
count = 0;
my_string = strtok(my_line, " \t\n");
while ( my_string != NULL ) {
    count++;
    printf("Thread %ld > string %d = %s\n", my_rank, count,
           my_string);
    my_string = strtok(NULL, " \t\n");
}

sem_wait(&sems[my_rank]);
fg_rv = fgets(my_line, MAX, stdin);
sem_post(&sems[next]);
}

return NULL;
} /* Tokenize */
```

# Running with one thread

- It correctly tokenizes the input stream.

Pease porridge hot.  
Pease porridge cold.  
Pease porridge in the pot  
Nine days old.

# Running with two threads

Thread 0 > my line = Pease porridge hot.

Thread 0 > string 1 = Pease

Thread 0 > string 2 = porridge

Thread 0 > string 3 = hot.

Thread 1 > my line = Pease porridge cold.

Thread 0 > my line = Pease porridge in the pot

Thread 0 > string 1 = Pease

Thread 0 > string 2 = porridge

Thread 0 > string 3 = in

Thread 0 > string 4 = the

Thread 0 > string 5 = pot

Thread 1 > string 1 = Pease

Thread 1 > my line = Nine days old.

Thread 1 > string 1 = Nine

Thread 1 > string 2 = days

Thread 1 > string 3 = old.

Oops!



# What happened?

- `strtok` caches the pointer to the input line by declaring a variable to have static storage class.
- This causes the value stored in this variable to persist from one call to the next.
- Unfortunately for us, this cached string is shared, not private.

# What happened?

- Thus, thread 0's call to `strtok` with the third line of the input has apparently overwritten the contents of thread 1's call with the second line.
- So the `strtok` function is not thread-safe. If multiple threads call it simultaneously, the output may not be correct.



# Other unsafe C library functions

- Regrettably, it's not uncommon for C library functions to fail to be thread-safe.
- The random number generator `random` in `stdlib.h`.
- The time conversion function `localtime` in `time.h`.
- In older version of POSIX systems `srand/rand` were also not thread safe

# "re-entrant" functions

- In some cases, the C standard specifies an alternate, thread-safe, version of a function.

```
char* strtok_r(  
    char*      string      /* in/out */,  
    const char* separators, /* in      */  
    char**     saveptr_p   /* in/out */);
```

- In principle "re-entrant" != "thread safe"
- In practice, re-entrant functions are often also thread safe
- When in doubt, check the documentation!
  - (man strtok\_r)

Questions?

Misc

# Static Initializers

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
```

```
pthread_rwlock_t rwlock = PTHREAD_RWLOCK_INITIALIZER;
```

# Thread Pinning

What if I want to force a thread to run on a specific core?

```
#define _GNU_SOURCE
#include <pthread.h>

void* thread_func(void* thread_args){
    ...
    cpu_set_t cpuset;
    pthread_t thread = pthread_self();
    /* Set affinity mask to include core 3 */
    CPU_ZERO(&cpuset);
    CPU_SET(3, &cpuset);
    s = pthread_setaffinity_np(thread, sizeof(cpuset), &cpuset);
    ...
}
```

# Timing Code

```
#include <sys/time.h>

#define GET_TIME(now) { \
    struct timeval t; \
    gettimeofday(&t, NULL); \
    now = t.tv_sec + t.tv_usec/1000000.0; \
}

int main(){
    ...
    double start, finish, elapsed;
    GET_TIME(start);
    ...
    // Code to be timed
    ...
    GET_TIME(finish);
    elapsed = finish - start;
    printf("The code to be timed took %e seconds\n", elapsed);
    ...
}
```

Questions?



# Concluding Remarks (1)

- A thread in shared-memory programming is analogous to a process in distributed memory programming.
- However, a thread is often lighter-weight than a full-fledged process.
- In Pthreads programs, all the threads have access to global variables, while local variables usually are private to the thread running the function.

# Concluding Remarks (2)

- When indeterminacy results from multiple threads attempting to access a shared resource such as a shared variable or a shared file, at least one of the accesses is an update, and the accesses can result in an error, we have a **race condition**.

# Concluding Remarks (3)

- A **critical section** is a block of code that updates a shared resource that can only be updated by one thread at a time.
- So the execution of code in a critical section should, effectively, be executed as serial code.

# Concluding Remarks (4)

- **Busy-waiting** can be used to avoid conflicting access to critical sections with a flag variable and a while-loop with an empty body.
- It can be very wasteful of CPU cycles.
- It can also be unreliable if compiler optimization is turned on.

# Concluding Remarks (5)

- A **mutex** can be used to avoid conflicting access to critical sections as well.
- Think of it as a lock on a critical section, since mutexes arrange for mutually exclusive access to a critical section.

# Concluding Remarks (6)

- A **semaphore** is the third way to avoid conflicting access to critical sections.
- It is an unsigned int together with two operations: `sem_wait` and `sem_post`.
- Semaphores are more powerful than mutexes since they can be initialized to any nonnegative value.

# Concluding Remarks (7)

- A **barrier** is a point in a program at which the threads block until all of the threads have reached it.
- A **read-write lock** is used when it's safe for multiple threads to simultaneously read a data structure, but if a thread needs to modify or write to the data structure, then only that thread can access the data structure during the modification.

# Concluding Remarks (8)

- Some C functions cache data between calls by declaring variables to be static, causing errors when multiple threads call the function.
- This type of function is not **thread-safe**.



# Chapter 5

## Shared Memory Programming with OpenMP

# Roadmap

- Writing programs that use OpenMP.
- Using OpenMP to parallelize many serial for loops with only small changes to the source code.
- Task parallelism.
- Explicit thread synchronization.
- Standard problems in shared-memory programming.

# OpenMP


- An API for shared-memory parallel programming.
- MP = multiprocessing
- Designed for shared-memory systems.
- System is viewed as a collection of cores or CPU's, all of which have access to main memory.

# OpenMP

- OpenMP aims to decompose a sequential program into components that can be executed in parallel.
- OpenMP allows an “incremental” conversion of sequential programs into parallel ones, with the assistance of the compiler.
- OpenMP relies on compiler directives for decorating portions of the code that the compiler will attempt to parallelize.

# OpenMP

- OpenMP programs are Globally Sequential, Locally Parallel.
- Programs follow the fork-join paradigm:



# Pragmas

- Special preprocessor instructions.
- Typically added to a system to allow behaviors that aren't part of the basic C specification.
- Compilers that don't support the pragmas ignore them.

`#pragma`

# OpenMP pragmas

- `# pragma omp parallel`
  - Most basic parallel directive.
  - The number of threads that run the following structured block of code is determined by the run-time system.

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
} /* Hello */
```



```
gcc -g -Wall -fopenmp -o omp_hello omp_hello.c
```

```
./omp_hello 4
```

compiling



running with 4 threads



```
Hello from thread 0 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 3 of 4
```

possible  
outcomes



```
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 0 of 4  
Hello from thread 3 of 4
```


```
Hello from thread 3 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 0 of 4
```

# Thread Team Size Control

- **Universally:** via the `OMP_NUM_THREADS` environmental variable:  
\$ echo \${OMP\_NUM\_THREADS} # to query the value  
\$ export OMP\_NUM\_THREADS=4 # to set it in BASH
- **Program level :** via the `omp_set_num_threads` function, outside an OpenMP construct.
- **Pragma level :** via the `num_threads` clause.
- **Precedence:**
  - Universally/env. Variable
  - Program level
  - Pragma level

# Thread Team Size Control

- The `omp_get_num_threads` call returns the active threads in a parallel region. If it is called in a sequential part it returns 1.
- The `omp_get_thread_num` returns the id of the calling thread (similar to the rank in MPI)



```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);


    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
} /* Hello */
```

# A process forking and joining two threads



# clause

- Text that modifies a directive.
- The `num_threads` clause can be added to a `parallel` directive.
- It allows the programmer to specify the number of threads that should execute the following block.

```
# pragma omp parallel num_threads ( thread_count )
```

# Of note...

- There may be system-defined limitations on the number of threads that a program can start.
- The OpenMP standard doesn't guarantee that this will actually start `thread_count` threads.
- Most current systems can start hundreds or even thousands of threads.
- Unless we're trying to start a lot of threads, we will almost always get the desired number of threads.
- After a block is completed, there is an implicit barrier.

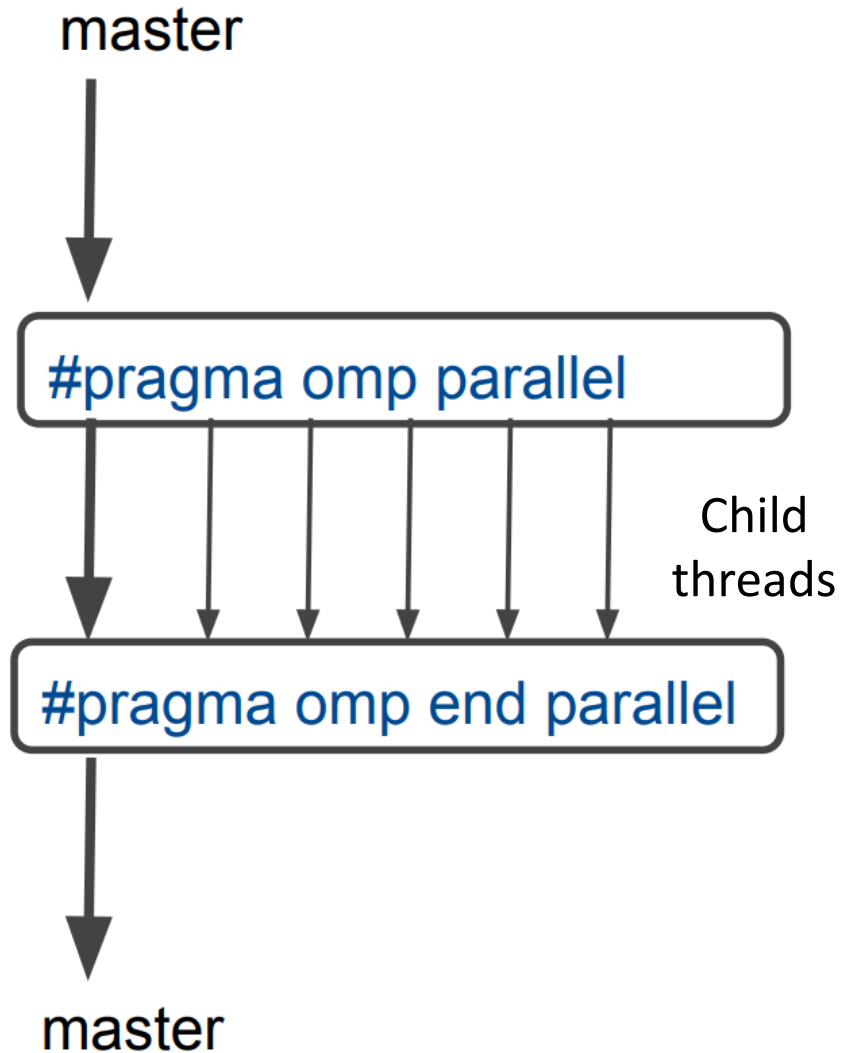
# Some terminology

- In OpenMP parlance the collection of threads executing the parallel block — the original thread and the new threads — is called a **team**
- **master:** the original thread of execution
- **parent:** thread that encountered a parallel directive and started a team of threads.
- In many cases, the parent is also the master thread.
- **child:** each thread started by the parent is considered a child thread.



# Parallel Construct

```
#pragma omp parallel  
{  
...  
}
```



# Parallel Construct

```
1  #include<stdio.h>
2  #include<omp.h>
3
4  int main()
5  {
6      printf("Only master thread here \n");
7
8      #pragma omp parallel
9      {
10         int tid = omp_get_thread_num();
11         printf("Hello I am thread number %d \n", tid);
12     }
13
14     printf("Only master thread here \n");
15 }
```

1

3

Only master thread here  
Hello I am thread number 0  
Hello I am thread number 1  
Hello I am thread number 2  
Hello I am thread number 3  
Only master thread here

2

int tid


4

*Each thread has a  
private copy of  
variable*

- Specifying number of threads
  - `export OMP_NUM_THREADS=4`
- Execute:
  - `./basic.exe`

# In case the compiler doesn't support OpenMP

```
# include <omp.h>
```



```
#ifdef _OPENMP  
# include <omp.h>  
#endif
```

# In case the compiler doesn't support OpenMP

```
# ifdef _OPENMP
    int my_rank = omp_get_thread_num ( );
    int thread_count = omp_get_num_threads ( );
# else
    int my_rank = 0;
    int thread_count = 1;
# endif
```