

Programmazione di Sistemi ~~Embedded e~~ Multicore

Teacher: Daniele De Sensi

Recap

Recap

- OpenMP basics
- OpenMP scope


Parallel For

Parallel for

- Forks a team of threads to execute the following structured block.
- However, the structured block following the parallel for directive must be a for loop.
- Furthermore, with the parallel for directive the system parallelizes the for loop by dividing the iterations of the loop among the threads.

Trapezoid Example

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```

Legal forms for parallelizable for statements

for	{	index = start ;		index++
				++index
			index < end	index--
			index <= end	--index
			index >= end ;	index += incr
			index > end	index -= incr
				index = index + incr
				index = incr + index
			index = index - incr	
)			

Why? It allows the runtime system to determine the number of iterations prior to the execution of the loop

Caveats

- The variable `index` must have integer or pointer type (e.g., it can't be a float).
- The expressions `start`, `end`, and `incr` must have a compatible type. For example, if `index` is a pointer, then `incr` must have integer type.
- The expressions `start`, `end`, and `incr` must not change during execution of the loop.
- During execution of the loop, the variable `index` can only be modified by the "increment expression" in the `for` statement.

examples

```
for (i=0; i<n; i++) {  
    if (...) break;    //cannot be parallelized  
}
```

```
for (i=0; i<n; i++) {  
    if (...) return 1; //cannot be parallelized  
}
```

```
for (i=0; i<n; i++) {  
    if (...) exit();   //can be parallelized  
}
```

```
for (i=0; i<n; i++) {  
    if (...) i++;      //CANNOT be parallelized  
}
```

Questions?

Example: Odd-Even Sort

Odd-Even Sort

This might fork/join new threads everytime it is called (depends on the implementation)

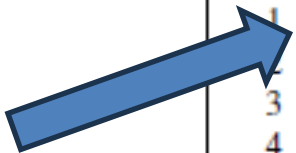
If it does so, we would have some **overhead**

Can we just create the threads at the beginning (before line 1)?

```
1   for (phase = 0; phase < n; phase++) {
2       if (phase % 2 == 0)
3       #   pragma omp parallel for num_threads(thread_count) \
4           default(none) shared(a, n) private(i, tmp)
5           for (i = 1; i < n; i += 2) {
6               if (a[i-1] > a[i]) {
7                   tmp = a[i-1];
8                   a[i-1] = a[i];
9                   a[i] = tmp;
10            }
11        }
12    else
13    #   pragma omp parallel for num_threads(thread_count) \
14        default(none) shared(a, n) private(i, tmp)
15        for (i = 1; i < n-1; i += 2) {
16            if (a[i] > a[i+1]) {
17                tmp = a[i+1];
18                a[i+1] = a[i];
19                a[i] = tmp;
20            }
21        }
22    }
```

Odd-Even Sort

Fork threads
only here



```
1  # pragma omp parallel num_threads(thread_count) \  
2    default(none) shared(a, n) private(i, tmp, phase)  
3    for (phase = 0; phase < n; phase++) {  
4      if (phase % 2 == 0)  
5        # pragma omp for  
6        for (i = 1; i < n; i += 2) {  
7          if (a[i-1] > a[i]) {  
8            tmp = a[i-1];  
9            a[i-1] = a[i];  
10           a[i] = tmp;  
11         }  
12       }  
13     else  
14       # pragma omp for  
15       for (i = 1; i < n-1; i += 2) {  
16         if (a[i] > a[i+1]) {  
17           tmp = a[i+1];  
18           a[i+1] = a[i];  
19           a[i] = tmp;  
20         }  
21       }  
22     }
```

Odd-Even Sort

Table 5.2 Odd-Even Sort with Two parallel for Directives and Two for Directives (times are in seconds)

thread_count	1	2	3	4
Two parallel for directives	0.770	0.453	0.358	0.305
Two for directives	0.732	0.376	0.294	0.239

Reusing the same threads provide faster execution times


Questions?

Nested Loops

Nested for loops

- If we have nested for loops, it is often enough to simply **parallelize the outermost loop**

```
a();  
#pragma omp parallel for  
for (int i = 0; i < 4; ++i) {  
    for (int j = 0; j < 4; ++j) {  
        c(i, j);  
    }  
}  
z();
```



Nested for loops

- Sometimes the outermost loop is so short that not all threads are utilized:
3 iterations, so it won't have sense to start more than 3 threads


```
a();  
#pragma omp parallel for  
for (int i = 0; i < 3; ++i) {  
    for (int j = 0; j < 6; ++j) {  
        c(i, j);  
    }  
}  
z();
```



Nested for loops

- We could try to parallelize the inner loop, but there is no guarantee that the thread utilization is better

```
a();  
for (int i = 0; i < 3; ++i) {  
    #pragma omp parallel for  
        for (int j = 0; j < 6; ++j) {  
            c(i, j);  
        }  
}  
z();
```



Nested for loops


- The correct solution is to **collapse it into one loop** that does 18 iterations. We can do it manually:

<https://ppc.cs.aalto.fi/ch3/nested/>

Nested for loops

- The correct solution is to **collapse it into one loop** that does 18 iterations. We can do it manually:


```
a();  
#pragma omp parallel for  
for (int ij = 0; ij < 3*6; ++ij) {  
    c(ij / 6, ij % 6);  
}  
z();
```



Nested for loops

- we can ask OpenMP to do it for us:


```
a();  
#pragma omp parallel for  
collapse(2)  
for (int i = 0; i < 3; ++i) {  
    for (int j = 0; j < 6; ++j)  
        c(i, j);  
}  
z();
```



Nested for loops

- **Wrong way:** “Nested parallelism” is disabled in OpenMP by default (i.e., inner parallel for pragmas will be ignored)

```
a();  
#pragma omp parallel for  
for (int i = 0; i < 3; ++i) {  
    #pragma omp parallel for  
        for (int j = 0; j < 6; ++j)  
            c(i, j);  
}  
z();
```



Nested for loops

- **Wrong way:** If "Nested parallelism" is enabled it will create 12 threads on a server with 4 cores (3*4)!


```
a();  
#pragma omp parallel for  
for (int i = 0; i < 3; ++i) {  
    #pragma omp parallel for  
        for (int j = 0; j < 6; ++j)  
            c(i, j);  
}  
z();
```


Data Dependencies


Chapter 4.4.1 from "Multicore and GPU
Programming An Integrated Approach"

Data dependencies

```
fibonacci[ 0 ] = fibonacci[ 1 ] = 1;  
for (i = 2; i < n; i++)  
    fibonacci[ i ] = fibonacci[ i - 1 ] + fibonacci[ i - 2 ];
```



```
fibonacci[ 0 ] = fibonacci[ 1 ] = 1;  
# pragma omp parallel for num_threads(2)  
  for (i = 2; i < n; i++)  
    fibonacci[ i ] = fibonacci[ i - 1 ] + fibonacci[ i - 2 ];
```



1 1 2 3 5 8 13 21 34 55

this is correct

but sometimes
we get this



1 1 2 3 5 8 0 0 0 0

What happened?

1. OpenMP compilers don't check for dependences among iterations in a loop that's being parallelized with a `parallel for` directive.
2. A loop in which the results of one or more iterations depend on other iterations cannot, in general, be correctly parallelized by OpenMP. We say that we have a **loop-carried dependence**

Data dependencies

- Assuming we have a loop of the form:

```
for ( i = ...  
{  
    S1 : operate on a memory location x  
    ...  
    S2 : operate on a memory location x  
}
```

- There are four different ways that S1 and S2 are connected, based on whether they are reading or writing to x.
- A problem exists if the dependence crosses loop iterations : **loop-carried dependence**.

Dependence Types

- Flow dependence : RAW

```
x = 10;          // S1  
y = 2 * x + 5;   // S2
```

- Anti-flow dependence : WAR

```
y = x + 3;       // S1  
x ++ ;           // S2
```

Dependence Types (cont.)

- Output dependence : WAW

```
x = 10;           // S1  
x = x + c;        // S2
```

- Input dependence : RAR (it's not an actual dependence)

```
y = x + c;        // S1  
z = 2 * x + 1;    // S2
```

Questions?

Flow Dependence Removal (RAW)

Data Dependency Resolution

6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

Data Dependency Resolution

6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

Flow Dependence : Reduction, Induction Variables

• Example:


```
double v = start;
double sum=0;
for (int i = 0; i < N; i++)
{
    sum = sum + f(v);    // S1
    v = v + step;        // S2
}
```

- RAW (S1) caused by reduction variable sum.
- RAW (S2) caused by induction variable v (induction variable is a variable that gets increased/decreased by a constant amount at each iteration).
- RAW (S2→S1) caused by induction variable v.
- Induction variable: affine function of the loop variable.

N.B.: RAW are between the i and the $i+1$ iterations. E.g., sum is **read** in the $(i+1)$ -th iteration **after** being **written** in i -th iteration

Remove RAW (S2) and RAW (S2->S1)

```
double v = start;
double sum = 0;
for(int i = 0; i < N ; i++)
{
    sum = sum + f(v); // S1
    v = v + step;      // S2
}
```



```
double v;
double sum = 0;
for(int i = 0; i < N ; i++)
{
    v = start + i*step;
    sum = sum + f(v);
}
```

i = 0 -> v = start
i = 1 -> v = start + step
i = 2 -> v = (start + step) + step
...

Remove RAW (S1)

```
double v;  
double sum = 0;  
#pragma omp parallel for reduction(+ : sum) private(v)  
for(int i = 0; i < N ; i++)  
{  
    v = start + i*step;  
    sum = sum + f(v);  
}
```

Data Dependency Resolution

6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

Flow Dependence: Loop Skewing


- Another technique involves the rearrangement of the loop body statements. Example with :

```
for (int i = 1; i < N; i++)  
{  
    y[ i ] = f( x[ i-1 ] ); // S1  
    x[ i ] = x[ i ] + c[ i ]; // S2  
}
```

- RAW (S2→S1) on x
- Solution: make sure the statements that consume the calculated values that cause the dependence, use values generated **during the same iteration.**

Flow Dependence : Loop Skewing (2)

```
for (int i = 1; i < N; i++)  
{  
    y[ i ] = f( x[ i-1 ] ); // S1  
    x[ i ] = x[ i ] + c[ i ]; // S2  
}
```



```
y[ 1 ] = f( x[ 0 ] );  
for (int i = 1; i < N - 1; i++)  
{  
    x[ i ] = x[ i ] + c[ i ];  
    y[ i + 1 ] = f( x[ i ] );  
}  
x[ N - 1 ] = x[ N - 1 ] + c[ N - 1 ];
```


Flow Dependence: Loop Skewing (3)

- How to do loop skewing?
 - Hint: unroll the loop and see the repetition pattern

```
for (int i = 1; i < N; i++)  
{  
    y[i] = f(x[i-1]);  
    x[i] = x[i] + c[i];  
}
```

```
y[1] = f(x[0]);  
x[1] = x[1] + c[1];  
y[2] = f(x[1]);  
x[2] = x[2] + c[2];  
  
...  
y[N-2] = f(x[N-3]);  
x[N-2] = x[N-2] + c[N-2];  
y[N-1] = f(x[N-2]);  
x[N-1] = x[N-1] + c[N-1];
```

Flow Dependence: Loop Skewing (4)

- How to do loop skewing?
 - Hint: unroll the loop and see the repetition pattern

```
for (int i = 1; i < N; i++)  
{  
    y[i] = f(x[i-1]);  
    x[i] = x[i] + c[i];  
}
```

```
y[1] = f(x[0]);  
x[1] = x[1] + c[1];  
y[2] = f(x[1]);  
x[2] = x[2] + c[2];  
  
...  
y[N-2] = f(x[N-3]);  
x[N-2] = x[N-2] + c[N-2];  
y[N-1] = f(x[N-2]);  
x[N-1] = x[N-1] + c[N-1];
```

Flow Dependence: Loop Skewing (5)

- How to do loop skewing?
 - Hint: unroll the loop and see the repetition pattern

```
y[1]= f(x[0]);  
for (int i = 1; i < N-1; i++)  
{  
    x[i]= x[i]+c[i];  
    y[i+1]= f(x[i]);  
}  
x[N-1]= x[N-1]+ c[N-1];
```

```
y[1]= f(x[0]);  
x[1]= x[1]+c[1];  
y[2]= f(x[1]);  
x[2]= x[2]+c[2];  
...  
y[N-2]= f(x[N-3]);  
x[N-2]= x[N-2]+ c[N-2];  
y[N-1]= f(x[N-2]);  
x[N-1]= x[N-1]+ c[N-1];
```

Data Dependency Resolution

6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

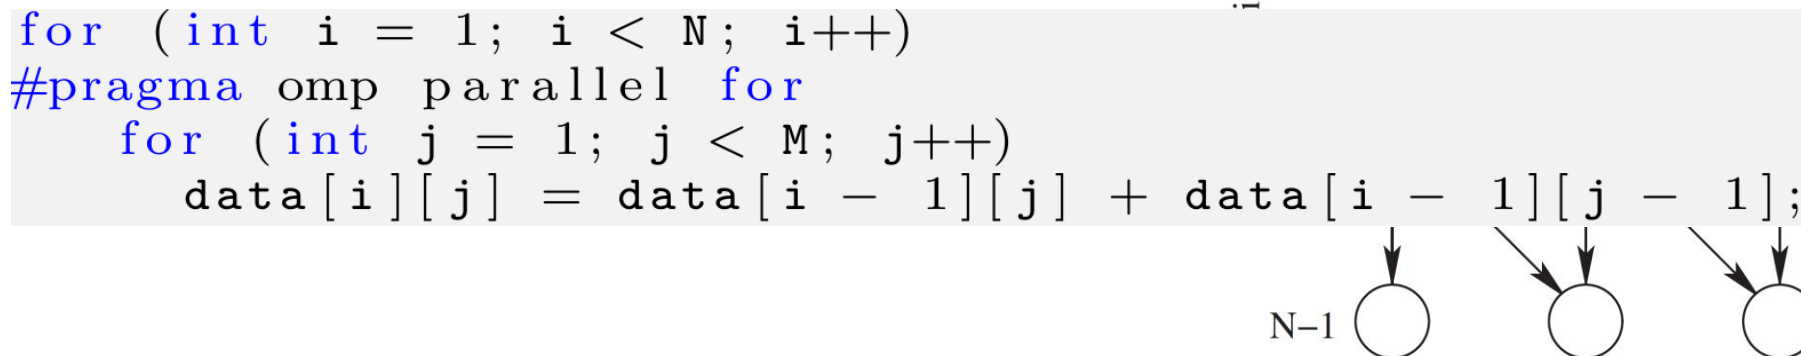
Iteration Space Dependency Graph

- ISDG is made up of nodes that represent an single execution of the loop body, and edges that represent dependencies.

- Example:

```
for (int i = 1; i < N; i++)
    for (int j = 1; j < M; j++)
        data[i][j] = data[i - 1][j] + data[i - 1][j - 1];
```

No edges/dependencies between nodes on the same row.
I.e., we can parallelize the j-loop



Data Dependency Resolution


6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

Flow Dependencies: Refactoring


- Refactoring refers to rewriting of the loop(s) so that parallelism can be exposed.
- The ISDG for the following example:

```
for (int i = 1; i < N; i++)  
  for (int j = 1; j < M; j++)  
    data[i][j] = data[i - 1][j] + data[i][j - 1] + data[i - 1][j - 1]; // S1
```




Flow Dependencies: Refactoring (2)

- Diagonal sets can be executed in parallel (no edges/dependencies between nodes in the same diagonal set):



Flow Dependencies: Refactoring (3)



```
for(wave=0 wave<NumWaves; wave++) {  
    diag=F(wave);  
    #pragma omp parallel for  
    for(k=0; k<diag; k++) {  
        int i = get_i(diag, k);  
        int j = get_j(diag, k);  
        data[i][j] = data[i-1][j] + data[i][j-1] + data[i-1][j-1];  
    }  
}
```

(Intuition, full code on the book)

The execution in waves requires a change of loop variables from the original i and j .

Data Dependency Resolution

6 techniques:


1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. **Fissioning**
6. Algorithm change

Flow Dependencies : Fissioning

Fissioning means breaking the loop apart into a sequential and a parallelizable part.

Example:

```
s = b[ 0 ];  
for (int i = 1; i < N; i++)  
{  
    a[ i ] = a[ i ] + a[ i - 1 ];    // S1  
    s = s + b[ i ];  
}
```



```
// sequential part  
for (int i = 1; i < N; i++)  
    a[ i ] = a[ i ] + a[ i - 1 ];  
  
// parallel part  
s = b[ 0 ];  
#pragma omp parallel for reduction(+ : s)  
for (int i = 1; i < N; i++)  
    s = s + b[ i ];
```

Data Dependency Resolution

6 techniques:

1. reduction/induction variable fix
2. Loop skewing
3. Partial parallelization
4. Refactoring
5. Fissioning
6. Algorithm change

Flow Dependencies: Algorithm Change

- If everything else fails, switching the algorithm maybe the answer.
- For example, the Fibonacci sequence:

```
for (int i = 2 ; i < N ; i++)  
{  
    int x = F[i - 2]; // S1  
    int y = F[i - 1]; // S2  
    F[i] = x + y;      // S3  
}
```

can be parallelized via Binet's formula:

$$F_n = \frac{\varphi^n - (1 - \varphi)^n}{\sqrt{5}}$$

Questions?

Antidependence Removal (WAR)

Antidependence Removal (WAR)

- E.g.:

```
for(int i = 0; i < N-1; i++)  
{  
    a[ i ] = a[ i + 1 ] + c;  
}
```
- Simple solution, make a copy of a before starting to modify it:

```
for(int i = 0; i < N-1; i++)  
{  
    a2[ i ] = a[ i + 1 ];  
}
```

```
#pragma omp parallel for  
for(int i = 0; i < N-1; i++)  
{  
    a[ i ] = a2[ i ] + c;  
}
```
- ATTENTION: Space and time tradeoffs must be carefully evaluated!

Output Dependence Removal (WAW)

Output Dependence Removal (WAW)

- E.g.:

```
for (int i = 0; i < N; i++)  
{  
    y[i] = a * x[i] + c; // S1  
    d = fabs( y[i] );    // S2  
}
```
- Where is the WAW dependency?
 - On variable d.
- How to guarantee that at the end of the execution the computed d is the one computed in the last iteration?

```
#pragma omp parallel for shared( a, c ) lastprivate( d )  
for (int i = 0; i < N; i++)  
{  
    y[i] = a * x[i] + c;  
    d = fabs( y[i] );  
}
```

Questions?

OPIS

OPIS

OPIS Code: B76Z7JS0

Instructions:

